

---

# Under the Shadow: Exploiting Opacity Variation for Fine-grained Shadow Detection

---

Xiaotian Qiao<sup>1,2</sup>, Ke Xu<sup>1</sup>, Xianglong Yang<sup>2</sup>, Ruijie Dong<sup>1</sup>, Xiaofang Xia<sup>1\*</sup>, Jiangtao Cui<sup>1\*</sup>

<sup>1</sup>School of Computer Science and Technology, Xidian University, China

<sup>2</sup>Guangzhou Institute of Technology, Xidian University, China

## Abstract

Shadow characteristics are of great importance for scene understanding. Existing works mainly consider shadow regions as binary masks, often leading to imprecise detection results and suboptimal performance for scene understanding. We demonstrate that such an assumption oversimplifies light-object interactions in the scene, as the scene details under either hard or soft shadows remain visible to a certain degree. Based on this insight, we aim to reformulate the shadow detection paradigm from the opacity perspective, and introduce a new fine-grained shadow detection method. In particular, given an input image, we first propose a shadow opacity augmentation module to generate realistic images with varied shadow opacities. We then introduce a shadow feature separation module to learn the shadow position and opacity representations separately, followed by an opacity mask prediction module that fuses these representations and predicts fine-grained shadow detection results. In addition, we construct a new dataset with opacity-annotated shadow masks across varied scenarios. Extensive experiments demonstrate that our method outperforms the baselines qualitatively and quantitatively, enhancing a wide range of applications, including shadow removal, shadow editing, and 3D reconstruction.

## 1 Introduction

Shadow is a natural phenomenon when objects obstruct light sources either fully or partially, resulting in intensity and color variations across certain regions. These shadow characteristics play vital roles for scene understanding. The presence of shadows can degrade image quality and impede numerous vision and graphics tasks, including image editing [48], object detection [33], and visual tracking. Furthermore, shadow analysis allows us to interpret semantics and geometry of a scene [31], *e.g.*, object shapes, light source positions, and object relationships. However, accurately detecting shadow characteristics (*e.g.*, boundary, opacity, etc.) in the scene can be very challenging. Shadows exhibit diverse characteristics due to scene factors like lighting conditions, surface properties, and occluder configurations, resulting in various shadow types, from sharp to diffuse.

While advanced shadow detection methods have been proposed, they often regard shadows as binary masks or assume uniform opacity within shadow regions [2, 17, 49, 42]. We argue that such a binary or simplified lighting paradigm oversimplifies light-object interactions, as the scene details under either hard or soft shadows retain varying degrees of visibility. As shown in Figure 1, existing methods [54, 55] fail to capture the nuanced opacity variations and suffer from position and shape estimation artifacts, undermining downstream applications like shadow removal and editing that rely on precise shadow characteristics. A fine-grained formulation of the shadow regions remains underexplored, which is vital for scene understanding.

---

\* Joint corresponding authors.

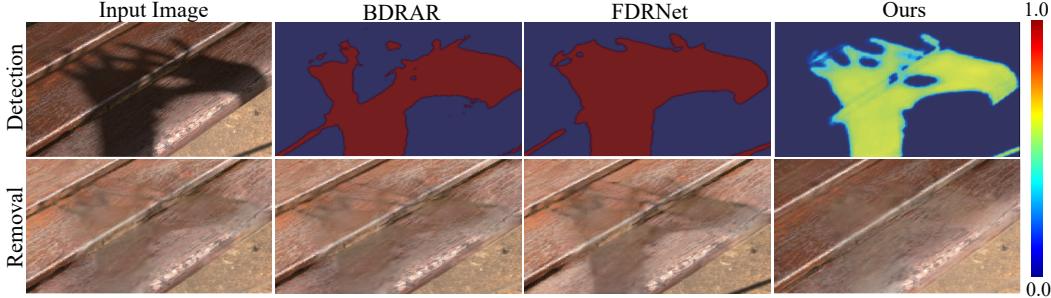


Figure 1: Fine-grained shadow detection. Given an input image (1st row), existing methods [17, 54, 55] predict binary shadow masks (1st, 2nd and 3rd columns), yielding inaccurate shadow detection results. In contrast, we propose to detect fine-grained shadow characteristics (4th column) that capture subtle opacity variations, with benefits for the downstream applications like shadow removal [12] (2nd row).

We observe that scene details under both umbra (fully occluded) and penumbra (partially occluded) regions retain varying degrees of visibility. The variation of shadow opacity caused by factors like occluder distance or light source intensity, could offer useful scene contextual cues. Take the input image in Figure 1 as an example. The subtle gradation of the shadow boundary suggests occluder distance, while umbra darkness and transparency reflect illumination intensity.

Inspired by the above observation, we propose a learning-based framework to predict a continuous map that indicates the position, boundary, and opacity variation of shadow regions. Our framework contains three main modules. First, given the input shadow image, we propose a Shadow Opacity Augmentation (SOA) module to generate multiple images with varied shadow opacities. Each augmented image and the original image are grouped as an image pair. Second, for each image pair, we propose a Shadow Feature Separation (SFS) module to disentangle position and opacity representations of the shadows separately. Third, we propose an Opacity Mask Prediction (OMP) module that fuses the learned representations and predicts a continuous shadow map as the output.

To support detailed analysis of the proposed paradigm, we further construct a new Fine-grained Shadow Detection (*i.e.*, FSD) dataset. In particular, it contains 2,653 images with opacity-annotated shadow masks across diverse scenarios, including different light source intensities and numbers, indoor and outdoor scenes. We conduct extensive qualitative and quantitative experiments on public datasets and the proposed dataset. Results show that our approach outperforms baselines in capturing fine-grained shadow characteristics, enhancing downstream scene understanding applications.

To sum up, we make the first attempt to investigate fine-grained shadow detection by exploiting opacity variations that are ignored by existing shadow detection and removal methods. We propose a new shadow detection method by explicitly capturing shadow position and opacity characteristics, and construct a new dataset with shadow opacity annotations across varied scenarios. Experimental results show that our method offers a promising shadow detection pipeline that can be applied across various applications, including shadow removal, shadow editing, and 3D reconstruction.

## 2 Related Work

**Shadow Detection.** Early research works rely on physical assumptions or heuristic cues. For instance, Lin et al. [29] adopted physical models assuming illumination invariance, while others analyzed chromatic variations in shadow regions using chromaticity [8, 9, 25], gradient [8, 25, 53], and intensity cues [39, 13, 53]. Huang et al. [20] trained a shadow detector by feeding edge features into a support vector machine. Guo et al. [14] computed illumination features of segmented regions and constructed a graph-based classifier by leveraging inter-region relationships. All these methods depend on handcrafted features, making it difficult to detect shadows accurately in complex scenes.

Deep learning has significantly improved the shadow detection performance [24]. Qu et al. [35] combined fully connected networks with patch-CNN refinement. Nguyen et al. [34] stabilized training using adversarial networks, and Hu et al. [17, 19, 54] leveraged spatial contexts. Wang et

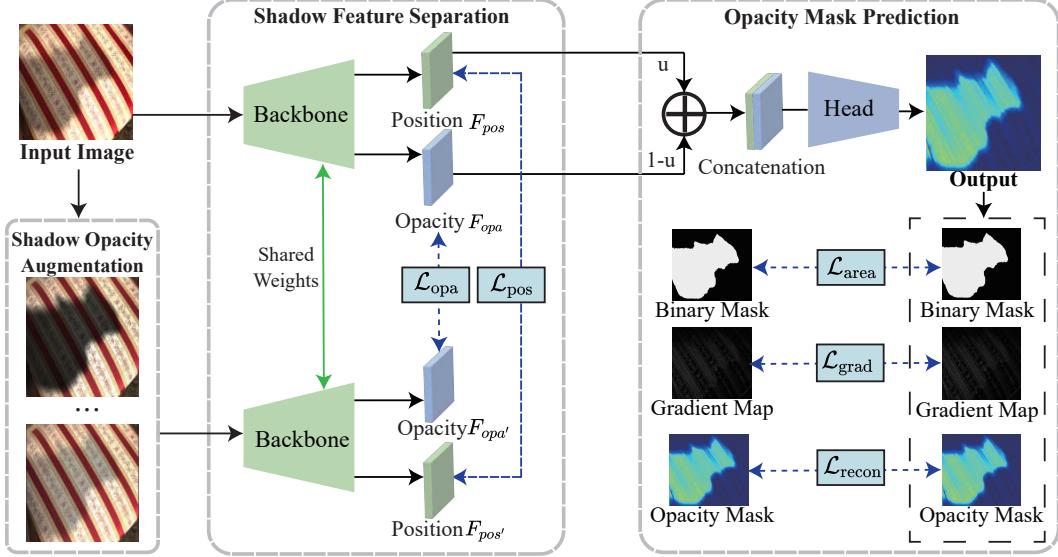


Figure 2: The overall pipeline of our fine-grained shadow detection.

al. [43, 6] and Zheng et al. [51] improved accuracy via adversarial networks and distraction modeling, respectively. Chen et al. [1] incorporated edge-assisted detection. Recent works [49, 36] further addressed annotation subjectivity and enhanced features via multi-color-space training.

However, these methods tend to produce binary masks, ignoring fine-grained shadow opacity. While Wang et al. [45] also noted the limitations of the binary shadow mask, they only generate soft masks implicitly as intermediate results for shadow removal. In contrast, our approach models the shadow opacity variations explicitly, benefiting various downstream scene understanding applications.

**Shadow Matte Estimation.** Shadow matte estimation aims to model the appearance of shadows in natural scenes. Chuang et al. [2] proposed matte shadows and defined a shadow compositing formula. Wu et al. [46] optimized the shadow compositing formulation and proposed a method to eliminate shadows using the shadow matte. Lu et al. [32] introduced a generic framework capable of matting objects with associated shadows. Lin et al. [28] further improved this to effectively matte objects and their corresponding shadows in 3D scenes.

However, shadow mattes involve additional colors and textures, making it difficult to be utilized directly for downstream applications. In contrast, our opacity-aware shadow mask formulation allows for more flexible and controllable shadow removal and editing results.

### 3 Approach

Given an input shadow image  $I_s$ , our goal is to predict a continuous opacity mask  $M_\alpha$  representing the shadow region. The shadow formulation is defined as follows:

$$I_s = I_f \cdot (1 - M_\alpha) + M_\alpha \cdot S, \quad (1)$$

where  $I_f$  is the shadow-free image.  $S$  represents the shadow color, which could be a constant value.  $M_\alpha$  represents the continuous shadow opacity, ranging from 0 (fully transparent) to 1(fully opaque).

Figure 2 shows the overall pipeline of our method, which contains a shadow opacity augmentation (SOA) module, a shadow feature separation (SFS) module, and an opacity mask prediction (OMP) module. We first apply opacity augmentation by adjusting the opacity of shadow regions in the input image. We then pass original and augmented images through the weights-shared backbone to obtain position and opacity features. The position features  $F_{pos}$  and  $F'_{pos}$  should be the same, while the opacity features  $F_{opa}$  and  $F'_{opa}$  should reflect the augmentation difference. Finally, both position and opacity features are fused through a shadow detection head to output the continuous opacity map.

### 3.1 Shadow Opacity Augmentation (SOA) Module

A trivial solution is to use an auto-encoder architecture to predict the opacity shadow mask directly. However, it performs poorly in handling diverse and complex shadow scenarios. To fully leverage the shadow opacity characteristics, we perform shadow augmentation on the input images by randomly augmenting the opacity of shadow regions.

Specifically, given an input image  $I_s$  and its opacity shadow mask  $M_\alpha$ , we modify the opacity of the shadow regions by randomly sampling the parameter  $\beta$  to generate an augmented shadow image  $\hat{I}_s$ . Note that the SOA module is only used for training, where  $I_s$  and  $\hat{I}_s$  are grouped as an image pair.

### 3.2 Shadow Feature Separation (SFS) Module

Given each image pair that indicates the same scene with different shadow opacities as input, we extract the position and opacity features separately. For the original image  $I_s$ , we introduce a backbone to extract the position feature  $F_{pos}$  and the opacity feature  $F_{opa}$ . Similarly, for the augmented shadow image  $\hat{I}_s$ , we use the weight-shared backbone to extract  $\hat{F}_{pos}$  and  $\hat{F}_{opa}$  correspondingly.

Since the shadow positions in the image pair remain unchanged, the position features  $F_{pos}$  and  $\hat{F}_{pos}$  should be identical theoretically. Meanwhile, the opacity features  $F_{opa}$  and  $\hat{F}_{opa}$  should exhibit the differences caused by the opacity augmentation parameter  $\beta$ .

### 3.3 Opacity Mask Prediction (OMP) Module

Given the set of the position feature  $F_{pos}$  and the opacity feature  $F_{opa}$ , we fuse these representations and predict the continuous shadow opacity map as the fine-grained shadow detection result. In particular, we first fuse the position feature  $F_{pos}$  and the opacity feature  $F_{opa}$  as follows:

$$F_s = \mu F_{pos} + (1 - \mu) F_{opa}. \quad (2)$$

We then employ a cumulative learning strategy [52] to optimize the fused shadow feature  $F_s$  through a weight selection strategy. For training epoch  $T$  and the total training epochs  $T_{max}$ ,  $\mu$  is defined as:

$$\mu = 1 - \left( \frac{T}{T_{max}} \right)^\eta, \quad (3)$$

where  $\eta$  is a hyperparameter controlling the decay rate of  $\mu$  during the training phase. Finally, the optimized shadow features are passed through a detection head to predict the opacity mask.

### 3.4 Training

For the SFS module, we use  $\mathcal{L}_{pos}$  to constrain the consistency of the position features:

$$\mathcal{L}_{pos} = \text{MAE}(F_{pos}, \hat{F}_{pos}), \quad (4)$$

where MAE is a mean-absolute-error loss function. The shadow opacity features  $F_{opa}$  and  $\hat{F}_{opa}$  should be capable of predicting the transparency augmentation parameter  $\beta$ . To achieve this, we subtract  $F_{opa}$  from  $\hat{F}_{opa}$  to predict  $\beta$ :

$$\mathcal{L}_{opa} = |\text{MLP}(F_{opa} - \hat{F}_{opa}) - \beta|, \quad (5)$$

where MLP contains a global average pooling layer and fully connected layers.

For the OMP module, to mitigate the computational difficulty arising from the varying proportions of shadow and non-shadow regions in each image, we employ  $\mathcal{L}_{area}$  to calculate the difference for shadow and non-shadow regions and multiply by their weights:

$$\mathcal{L}_{area} = \frac{\sum_i L_s \left( A_i^{\text{pred}}, A_i^{\text{gt}} \right) \cdot [A_i^{\text{gt}} = 0]}{\sum_i [A_i^{\text{gt}} = 0]} + \frac{\sum_i L_s \left( A_i^{\text{pred}}, A_i^{\text{gt}} \right) \cdot [A_i^{\text{gt}} \neq 0]}{\sum_i [A_i^{\text{gt}} \neq 0]}, \quad (6)$$




Figure 3: Example images with opacity mask annotations in our FSD dataset.

where  $A_i^{\text{pred}}$  represents the predicted value of the  $i$ -th pixel, and  $A_i^{\text{gt}}$  represents the true value of the  $i$ -th pixel. The definition of  $L_s$  is:

$$L_s(\text{pred}, \text{gt}) = \begin{cases} 0.5 \cdot (\text{pred} - \text{gt})^2 & \text{if } |\text{pred} - \text{gt}| < 1 \\ |\text{pred} - \text{gt}| - 0.5 & \text{otherwise.} \end{cases} \quad (7)$$

To achieve smooth prediction results that are close to the real shadow values, we also use gradient loss  $\mathcal{L}_{\text{grad}}$  as follows:

$$\mathcal{L}_{\text{grad}} = (\|\nabla_x \text{pred} - \nabla_x \text{gt}\| + \|\nabla_y \text{pred} - \nabla_y \text{gt}\|), \quad (8)$$

where  $\nabla_x$  and  $\nabla_y$  represent the gradients in the x and y directions, respectively.

We further utilize  $\mathcal{L}_{\text{recon}}$  to address the overall discrepancy as follows:

$$\mathcal{L}_{\text{recon}} = \frac{1}{N} \sum_{i=1}^N (\text{pred} - \text{gt})^2. \quad (9)$$

In summary, we train our model with the following losses:

$$\mathcal{L}_{\text{all}} = \lambda_{\text{pos}} \mathcal{L}_{\text{pos}} + \lambda_{\text{opa}} \mathcal{L}_{\text{opa}} + \lambda_{\text{area}} \mathcal{L}_{\text{area}} + \lambda_{\text{grad}} \mathcal{L}_{\text{grad}} + \lambda_{\text{recon}} \mathcal{L}_{\text{recon}}, \quad (10)$$

where  $\lambda_{\text{pos}}$ ,  $\lambda_{\text{opa}}$ ,  $\lambda_{\text{area}}$ ,  $\lambda_{\text{grad}}$ , and  $\lambda_{\text{recon}}$  are the weighting parameters.

## 4 Dataset

A number of shadow datasets [35, 43, 18] have been proposed in the community, advancing the progress of shadow detection and removal tasks. However, as shown in Table 1, these datasets suffer from shortcomings like lacking scenes with weak lighting, non-point light sources, or multi-source soft shadows. Most importantly, there are no publicly available shadow datasets considering opacity characteristics.

Dataset	Opacity Mask	Binary Mask	Object Association	Shadow Free
SRD [35]	✗	✗	✗	✓
SBU [41]	✗	✓	✗	✗
ISTD [43]	✗	✓	✗	✓
USR [30]	✗	✗	✗	✓
SOBA [44]	✗	✓	✓	✗
DESOBA [16]	✗	✓	✓	✓
RdSOBA [38]	✗	✓	✓	✓
FSD (ours)	✓	✓	✓	✓

Table 1: A taxonomy of shadow datasets.

To increase the diversity of the images, we use Blender to construct a new synthetic fine-grained shadow detection (i.e., FSD) dataset with the following guidelines:

- **Object.** We select common and diverse object categories from daily life (e.g., person, animals, shoes), and place 1 to 5 objects with randomized positions in each scene.
- **Scene.** The shadow images were captured from different camera angles across various indoor and outdoor scenarios, with all cameras configured at a 50mm focal length, automatic sensor settings, and randomized positions.
- **Light Source.** To enhance diversity, the number, position, and intensity of light sources are considered, with their positions and brightness levels being randomly sampled.

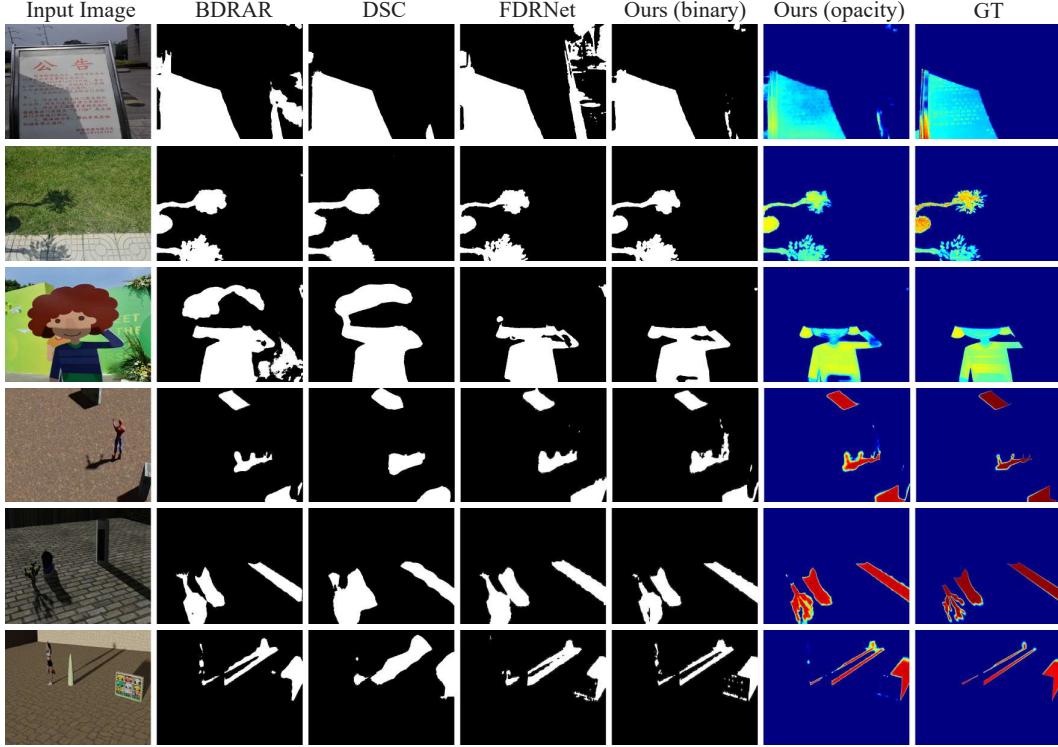


Figure 4: Qualitative comparisons of shadow detection on the ISTD [43] and FSD datasets. Given the input image (1st column), we show the baseline results (2nd to 4th columns), our binary mask(5th column), our opacity mask (6th column), and ground truth (7th column), respectively.

In summary, FSD comprises 2,653 scenes with different object, scene, and light source properties. Each scene contains varied camera positions, light intensities, object numbers, and categories, resulting in 7,701 object-shadow pairs. Some examples with varying shadow opacities are shown in Figure 3. Please refer to the supplementary material for additional details.

## 5 Experiments

### 5.1 Experimental Setup

**Implementation Details.** All experiments are conducted on a single NVIDIA RTX 4090 GPU. We use EfficientNet-B3 [37] as the backbone network and ImageNet [5] for weights initialization. Our model is trained using the Adam optimizer for 20 epochs, with an initial learning rate of 5e-4, dynamically adjusted through an exponential decay strategy (decay rate of 0.7). During training, the input images are resized to  $400 \times 400$  with a batch size of 12. The random horizontal flipping operation is further used for data augmentation. The weight parameters  $\lambda_{pos}$ ,  $\lambda_{opa}$ ,  $\lambda_{area}$ ,  $\lambda_{grad}$ , and  $\lambda_{recon}$  are set to 1, 1, 10, 10, and 10, respectively.

**Datasets.** We conduct experiments on both the proposed FSD dataset and the existing ISTD [43] dataset, which contains 1,330 training images and 540 test images. Furthermore, we utilize two shadow removal datasets, *i.e.*, ISTD+ [26] and SRD [35], for downstream applications. The ground truth shadow masks in SRD are obtained by utilizing a connectivity-based optimization strategy.

**Compared Methods.** As this is the first work for fine-grained shadow detection that considers opacity, to make fair comparisons, we convert our output to binary shadow masks, and compare our method with existing shadow detection methods, including SDCM [56], FDRNet [55], MTMT [1], DSD [51], DSC [17], BDRAR [54], ST-CGAN [43], scGAN [34], stacked-CNN [41].

Dataset	Method	BER ↓	Shadow ↓	Non Shad.↓
ISTD	stacked-CNN [41]	8.60	7.69	9.23
	scGAN [34]	4.70	3.22	6.18
	ST-CGAN [43]	3.85	2.14	5.55
	BDRAR [54]	2.69	<b>0.50</b>	4.87
	DSC [17]	3.42	3.85	3.00
	DSD [51]	2.17	1.36	2.98
	MTMT [1]	1.72	1.36	2.08
	FDRNet [55]	1.55	1.22	1.88
	SDCM [56]	<u>1.40</u>	1.02	<u>1.78</u>
	Ours	<b>1.32</b>	<u>0.96</u>	<b>1.67</b>
FSD	DSC [17]	4.46	4.99	3.93
	BDRAR [54]	4.43	7.74	<b>1.12</b>
	ST-CGAN [43]	4.19	4.63	3.76
	FDRNet [55]	3.01	<u>2.58</u>	3.43
	Ours	<b>1.79</b>	<b>1.94</b>	<u>1.64</u>

Table 2: Quantitative comparison of the proposed method with baselines on the ISTD [43] and FSD datasets. The best and second-best performances are marked in bold and underlined.

**Evaluation Metrics.** We evaluate the quality of the predicted fine-grained shadow mask from different aspects. We use the Balanced Error Rate (BER) metric to measure the position and boundary of the shadows. In addition, to evaluate the accuracy of predicted shadow opacity, we propose a new metric termed Weighted Shadow Error (WSE). Since non-shadow areas typically dominate most images, commonly used metrics like RMSE may overlook subtle opacity variations within shadow regions. As a result, WSE incorporates a weighted strategy that emphasizes shadow region accuracy by computing a region-aware RMSE, thereby offering a more balanced assessment of fine-grained shadow detection performance.

$$WSE = \left( \frac{1 - \frac{\sum_{i=1}^N [A_i^{gt} \neq 0]}{N}}{\frac{\sum_{i=1}^N [A_i^{gt} \neq 0]}{N}} \right) \cdot RMSE_{sdr} + RMSE_{nsdr}, \quad (11)$$

where  $A_i^{gt}$  represents the alpha value of shadow for each pixel.  $RMSE_{sdr}$  and  $RMSE_{nsdr}$  are the root mean square error of the shadow and non-shadow regions, respectively.

## 5.2 Results

**Qualitative Evaluation.** Visual comparisons on the ISTD [43] and FSD datasets are presented in Figure 4. Different from existing methods that predict only binary masks for shadow regions, our approach can further capture the degree of shadow degradation in the scene accurately. Moreover, by explicitly modeling shadow opacity, our method substantially reduces both false positives and false negatives. Take the second row as an example. Our method delivers fine-grained detection results around the shadow boundaries of the tree leaves, demonstrating clear advantages in predicting precise shadow characteristics.

**Quantitative Evaluation.** We compare the performance of our method with state-of-the-art shadow detection methods in Table 2. As the existing shadow datasets do not contain ground truth opacity annotations, we extract the brightness difference between shadowed and shadow-free images as the ground truth value for the opacity shadow mask. Our method achieves the best or second-best BER scores among all compared methods, significantly validating the effectiveness of the shadow opacity guidance. Note that our method achieves a greater performance gain on the FSD dataset than on the ISTD dataset. The primary reason is that our FSD dataset includes a diverse range of challenging cases featuring soft shadow boundaries and varied opacity levels. Thus, these results highlight the strength of our method in handling diverse shadow characteristics in complex scenarios.

Method	BER ↓	RMSE ↓	WSE ↓
AutoEncoder	4.01	0.0919	10.49
w/o $F_{opa}$	<u>1.97</u>	-	-
w/o $F_{pos}$	2.82	<b>0.0637</b>	<u>5.93</u>
w/o SOA	2.16	0.0744	6.02
Ours (full)	<b>1.78</b>	<u>0.0699</u>	<b>5.78</b>

Table 3: Ablation study of different components.

Method	BER ↓	RMSE ↓	WSE ↓
w/o $\mathcal{L}_{grad}$ and $\mathcal{L}_{opa}$	3.27	-	-
w/o $\mathcal{L}_{grad}$ and $\mathcal{L}_{area}$	2.34	0.0725	6.81
w/o $\mathcal{L}_{grad}$	2.03	0.0691	6.74
w/o $\mathcal{L}_{area}$	<u>1.86</u>	<u>0.0682</u>	<u>6.03</u>
Ours (full)	<b>1.78</b>	<b>0.0669</b>	<b>5.78</b>

Table 4: Ablation study of different losses.

### 5.3 Ablation Study

We evaluate the impact of different components, losses, and training strategies of our pipeline on the FSD dataset. Further analysis on different datasets are provided in the supplementary material.

**Component Analysis.** We first introduce a naive baseline (i.e., AutoEncoder) with a simple encoder and decoder architecture. We then conduct two ablations for the SFS module. Specifically, we remove the position feature branch (i.e., w/o  $F_{pos}$ ) and the opacity feature branch (i.e., w/o  $F_{opa}$ ) individually. We further remove the SOA module (i.e., w/o SOA) to examine the effect of the shadow opacity augmentation strategy. As shown in Table 3, these components play essential roles in resolving ambiguities in fine-grained shadow characteristics.

Method	BER ↓	RMSE ↓	WSE ↓
w/o cum.	2.18	0.1505	10.83
$\eta = 0.3$	1.93	<b>0.0635</b>	6.21
$\eta = 0.4$	1.83	0.0984	<u>6.16</u>
$\eta = 0.6$	<u>1.80</u>	0.0675	6.17
$\eta = 0.5$	<b>1.78</b>	<u>0.0669</u>	<b>5.78</b>

Table 5: Ablation study of the training strategy.

**Loss Analysis.** The effects of different losses are examined in Table 4. We first remove  $\mathcal{L}_{grad}$  and  $\mathcal{L}_{opa}$  to study the importance of shadow opacity guidance. We then remove  $\mathcal{L}_{grad}$  and  $\mathcal{L}_{area}$  separately or totally. The results show that compared to using  $\mathcal{L}_{recon}$  alone, the integration of  $\mathcal{L}_{grad}$ ,  $\mathcal{L}_{area}$ , and  $\mathcal{L}_{opa}$  enhance the fine-grained shadow detection performance significantly.

**Training Analysis.** We also ablate the cumulative learning strategy in Table 5. We retrain the model without using the cumulative learning strategy (i.e., w/o cum.), and use different  $\eta$  values to control the decay rate of  $\mu$ . The results show that  $\eta = 0.5$  achieves the best performance across all metrics.

	Methods	Input Masks	All			Shadow			Non-Shadow		
			PSNR↑	SSIM↑	MAE↓	PSNR↑	SSIM↑	MAE↓	PSNR↑	SSIM↑	MAE↓
ISTD+	DC-ShadowNet [22]	NA	25.03	0.926	7.77	31.06	0.976	12.62	27.03	0.961	6.82
	DeS3 [23]	NA	31.38	0.958	3.94	36.49	0.989	6.56	34.70	0.972	3.40
	BMNet [57]	FDRNet [55]	32.41	0.962	3.40	36.59	0.989	6.65	35.96	0.978	2.83
	HomoFormer [47]	FDRNet [55]	32.41	0.953	3.51	38.84	0.991	5.31	34.58	0.966	3.17
	ShadowDiffusion [12]	FDRNet [55]	<b>34.08</b>	<u>0.968</u>	<u>3.12</u>	<b>40.12</b>	<b>0.992</b>	<b>5.15</b>	<u>36.66</u>	<u>0.978</u>	<u>2.74</u>
	ShadowDiffusion [12]	Ours	34.06	<b>0.969</b>	<b>3.11</b>	39.49	<b>0.992</b>	<u>5.28</u>	<b>37.22</b>	<b>0.981</b>	<b>2.68</b>
SRD	DC-ShadowNet [22]	NA	31.53	0.955	4.65	34.00	0.975	7.70	35.53	0.981	3.65
	DeS3 [23]	NA	34.11	0.968	3.56	37.91	0.986	5.27	37.45	0.984	3.03
	SAM-helps-shadow [50]	NA	30.72	0.952	4.79	33.94	0.979	7.44	33.85	0.981	3.74
	ShadowFormer [11]	DHAN [3]	32.46	0.957	4.28	35.55	0.982	6.14	36.82	0.983	3.54
	DHAN [3]	DHAN [3]	30.51	0.949	5.67	33.67	0.978	8.94	34.79	0.979	4.80
	BMNet [57]	DHAN [3]	31.69	0.956	4.46	35.05	0.981	6.61	36.02	0.982	3.61
	Inpaint4Shadow [27]	DHAN [3]	33.27	0.967	3.81	36.73	0.985	5.70	36.70	0.985	3.27
	Homoformer [47]	DHAN [3]	<b>35.37</b>	<u>0.972</u>	<u>3.33</u>	38.81	0.987	<u>4.25</u>	<b>39.45</b>	<b>0.988</b>	<b>2.85</b>
	ShadowDiffusion [12]	DHAN [3]	34.73	0.970	3.63	38.72	0.987	4.98	37.78	0.985	3.44
	ShadowDiffusion [12]	FDRNet [55]	34.23	<u>0.972</u>	3.50	37.31	0.985	5.04	<u>38.61</u>	<u>0.986</u>	<u>2.90</u>
	ShadowDiffusion [12]	Ours	<u>34.84</u>	<b>0.974</b>	<u>3.27</u>	<b>42.06</b>	<b>0.995</b>	<u>3.32</u>	36.45	<u>0.986</u>	3.28

Table 6: Quantitative comparisons with the shadow removal methods on ISTD+ [26] and SRD [35] datasets. NA indicates that no mask input is required. The best and the second results are highlighted in bold and underlined, respectively.



Figure 5: Shadow removal results. Given the input image (1st row), we show shadow removal results from ShadowDiffusion [12] with binary masks (2nd row), our opacity masks (3rd row), and the ground truth (4th row).



Figure 6: Shadow editing results. Given the input image (1st column), we extract the opacity shadow mask with the corresponding object (2nd column). Multiple editing operations can be further conducted, and corresponding fine-grained shadow characteristics will be changed adaptively. For example, replacing the background (3rd column) or changing the foreground object (4th column).

## 5.4 Applications

Benefiting from our model for fine-grained shadow detection, we explore three downstream applications here. Please refer to the supplementary material for additional details and results.

**Shadow Removal.** The opacity shadow mask provides rich information by explicitly encoding pixel-wise attenuation coefficients, enabling precise characterization of shadow degradation levels on the background. To validate its practical utility, we reformulate the training pipeline of ShadowDiffusion [12] by replacing its binary shadow mask input with our continuous opacity map. Such an adaptation can dynamically adjust denoising strength based on shadow opacity and preserve natural illumination transitions at shadow boundaries through opacity-guided attention modulation.

Quantitative comparisons against existing shadow removal methods are presented in Table 6. The results show that the shadow removal model trained with the opacity shadow mask achieves substantial improvements. Note that our method obtains greater performance gains on shadow regions than on non-shadow regions. The primary reason is that our method is specifically designed to handle shadow opacity, while non-shadow regions dominate most images, masking shadow-specific improvements. For example, on the SRD dataset [35] that contains more complex scenes and soft shadows, the PSNR value for shadow regions shows substantial enhancement, validating the effectiveness of the opacity guidance strategy. Qualitative comparisons in Figure 5 further highlight our advantages in shadow removal visual quality, particularly for shadows with light-color or blurred boundaries.

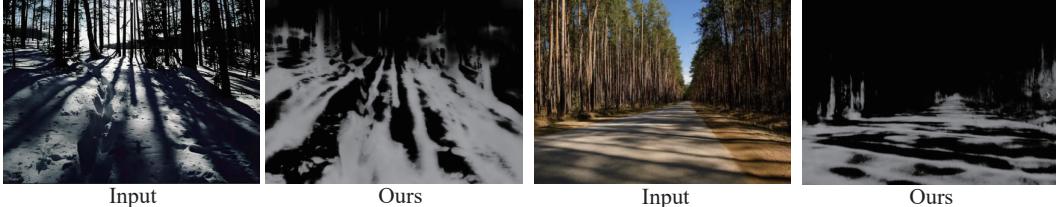


Figure 7: Failure cases. Our model may fail on images with complex shadow interactions.

**Shadow Editing.** Shadows often complicate image editing by hindering consistent manipulation of objects and their associated shading. However, by employing the proposed opacity shadow mask, we can simplify such a process by enabling synchronized adjustments of shadows and objects. As shown in Figure 6, the opacity shadow mask can be treated as a detachable shadow layer and edited alongside the object. Specifically, given an input image, we utilize the proposed model to extract an opacity shadow mask with the corresponding object. Multiple editing operations can be further conducted, and corresponding fine-grained shadow characteristics will be changed adaptively. For example, we can replace the background in the input image with a new one, or modify the position and size of the corresponding object in the image.

**3D reconstruction.** Fine-grained shadow opacity variations (e.g., penumbra width and transparency gradients) provide critical cues about scene geometry, such as occluder distance and light source intensity, which are highly valuable for 3D reconstruction. Following the shadow-driven neural rendering framework [40], we apply our continuous opacity maps, rather than binary masks from FDRNet [55] used before, for the 3D reconstruction application. Our method achieved RMSE values of 0.00957 and 0.00531 for the Bunny and CUBE scenes, respectively, significantly lower than FDRNet [55]’s 0.01074 and 0.00777. These results demonstrate that leveraging fine-grained shadow properties can enhance geometric reconstruction accuracy, particularly under texture-less or complex lighting conditions.

**Discussion.** The transparency information embedded in the opacity shadow mask provides additional benefits for scene understanding. First, it enables precise modeling of shadow-object interactions by quantifying the opacity gradient at shadow boundaries, which can guide existing shadow generation models to synthesize more realistic soft shadows with adaptive opacity. The pixel-wise opacity map serves as a physical prior for illumination decomposition, allowing joint optimization of shadow removal and scene relighting tasks. In dynamic scenes with moving objects, the temporal consistency of opacity values can be leveraged to enhance video shadow stabilization algorithms, reducing flickering artifacts during editing.

Although our model works well in various shadow scenarios, it may fail in some challenging cases. As illustrated in Figure 7, it would be difficult for our model to predict accurate shadow opacity when multiple shadow interactions exist in complex scenarios. A potential solution to this problem is to adopt the instance shadow strategy to distinguish specific object-shadow associations. As future work, we would like to explore more shadow characteristics to boost the scene understanding ability.

## 6 Conclusion

In this paper, we make the first attempt to investigate fine-grained shadow detection by exploiting shadow opacity characteristics in the scene. To this end, we propose a learning-based model that explicitly captures shadow position and opacity variations. In addition, we construct the first fine-grained shadow detection (FSD) dataset with opacity annotations across varied scenarios. Extensive qualitative and quantitative results show that our model can predict fine-grained shadow characteristics, achieving superior performance over the baselines and enhancing a wide range of applications, including shadow removal, shadow editing, and 3D reconstruction.

## Acknowledgments

This work was partially supported by Guangdong Basic and Applied Basic Research Foundation (No.2022A1515110740), National Natural Science Foundation of China (No.62302356, No.62372360, No.62372352), CCF-ALIMAMA TECH Kangaroo Fund (NO.CCF-ALIMAMA OF 2025007), Youth Elite Scientist Sponsorship Program by China Association for Science and Technology (No.YESS20220610), and Innovation Capability Support Program of Shaanxi Province (No.2024ZC-KJXX-021).

## References

- [1] Zhihao Chen, Lei Zhu, Liang Wan, Song Wang, Wei Feng, and Pheng-Ann Heng. A multi-task mean teacher for semi-supervised shadow detection. In *CVPR*, 2020.
- [2] Yung-Yu Chuang, Dan B Goldman, Brian Curless, David H Salesin, and Richard Szeliski. Shadow matting and compositing. In *SIGGRAPH*, 2003.
- [3] Xiaodong Cun, Chi-Man Pun, and Cheng Shi. Towards ghost-free shadow removal via dual hierarchical aggregation network and shadow matting gan. In *AAAI*, 2020.
- [4] Lennart Demes. ambientcg - cc0 textures, hdris and models, July 2024. URL <https://ambientcg.com/>.
- [5] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale hierarchical image database. In *CVPR*, 2009.
- [6] Bin Ding, Chengjiang Long, Ling Zhang, and Chunxia Xiao. Argan: Attentive recurrent generative adversarial network for shadow detection and removal. In *ICCV*, 2019.
- [7] Laura Downs, Anthony Francis, Nate Koenig, Brandon Kinman, Ryan Hickman, Krista Reymann, Thomas B. McHugh, and Vincent Vanhoucke. Google scanned objects: A high-quality dataset of 3d scanned household items, 2022. URL <https://arxiv.org/abs/2204.11918>.
- [8] Graham D Finlayson, Steven D Hordley, Cheng Lu, and Mark S Drew. On the removal of shadows from images. *TPAMI*, 28(1):59–68, 2005.
- [9] Graham D Finlayson, Mark S Drew, and Cheng Lu. Entropy minimization for shadow removal. *IJCV*, 2009.
- [10] Maciej Gryka, Michael Terry, and Gabriel J Brostow. Learning to remove soft shadows. *ACM TOG*, 2015.
- [11] Lanqing Guo, Siyu Huang, Ding Liu, Hao Cheng, and Bihan Wen. Shadowformer: Global context helps image shadow removal. In *AAAI*, 2023.
- [12] Lanqing Guo, Chong Wang, Wenhan Yang, Siyu Huang, Yufei Wang, Hanspeter Pfister, and Bihan Wen. Shadowdiffusion: When degradation prior meets diffusion model for shadow removal. In *CVPR*, 2023.
- [13] Ruiqi Guo, Qieyun Dai, and Derek Hoiem. Single-image shadow detection and removal using paired regions. In *CVPR*, 2011.
- [14] Ruiqi Guo, Qieyun Dai, and Derek Hoiem. Paired regions for shadow detection and removal. *IEEE TPAMI*, 2012.
- [15] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Delving deep into rectifiers: Surpassing human-level performance on imagenet classification. In *ICCV*, pages 1026–1034, 2015.
- [16] Yan Hong, Li Niu, and Jianfu Zhang. Shadow generation for composite image in real-world scenes. In *AAAI*, 2022.
- [17] Xiaowei Hu, Chi-Wing Fu, Lei Zhu, Jing Qin, and Pheng-Ann Heng. Direction-aware spatial context features for shadow detection and removal. *IEEE TPAMI*, 2019.

- [18] Xiaowei Hu, Yitong Jiang, Chi-Wing Fu, and Pheng-Ann Heng. Mask-shadowgan: Learning to remove shadows from unpaired data. In *ICCV*, 2019.
- [19] Xiaowei Hu, Tianyu Wang, Chi-Wing Fu, Yitong Jiang, Qiong Wang, and Pheng-Ann Heng. Revisiting shadow detection: A new benchmark dataset for complex world. *IEEE TIP*, 2021.
- [20] Xiang Huang, Gang Hua, Jack Tumblin, and Lance Williams. What characterizes a shadow boundary under the sun and sky? In *ICCV*, 2011.
- [21] Sketchfab Inc. Sketchfab - the best 3d viewer on the web, July 2024. URL <https://sketchfab.com/>.
- [22] Yeying Jin, Aashish Sharma, and Robby T Tan. Dc-shadownet: Single-image hard and soft shadow removal using unsupervised domain-classifier guided network. In *ICCV*, 2021.
- [23] Yeying Jin, Wei Ye, Wenhan Yang, Yuan Yuan, and Robby T Tan. Des3: Adaptive attention-driven self and soft shadow removal using vit similarity. In *AAAI*, 2024.
- [24] Salman H Khan, Mohammed Bennamoun, Ferdous Sohel, and Roberto Togneri. Automatic shadow detection and removal from a single image. *IEEE TPAMI*, 2015.
- [25] Jean-François Lalonde, Alexei A Efros, and Srinivasa G Narasimhan. Detecting ground shadows in outdoor consumer photographs. In *ECCV*, 2010.
- [26] Hieu Le and Dimitris Samaras. Shadow removal via shadow image decomposition. In *ICCV*, 2019.
- [27] Xiaoguang Li, Qing Guo, Rabab Abdelfattah, Di Lin, Wei Feng, Ivor Tsang, and Song Wang. Leveraging inpainting for single-image shadow removal. In *ICCV*, 2023.
- [28] Geng Lin, Chen Gao, Jia-Bin Huang, Changil Kim, Yipeng Wang, Matthias Zwicker, and Ayush Saraf. Omnimatterf: Robust omnimatte with 3d background modeling. In *ICCV*, 2023.
- [29] Yun-Hsuan Lin, Wen-Chin Chen, and Yung-Yu Chuang. Bedsr-net: A deep shadow removal network from a single document image. In *CVPR*, 2020.
- [30] Daquan Liu, Chengjiang Long, Hongpan Zhang, Hanning Yu, Xinzhi Dong, and Chunxia Xiao. Arshadowgan: Shadow generative adversarial network for augmented reality in single light scenes. In *CVPR*, 2020.
- [31] Ruoshi Liu, Sachit Menon, Chengzhi Mao, Dennis Park, Simon Stent, and Carl Vondrick. What you can reconstruct from a shadow. In *CVPR*, 2023.
- [32] Erika Lu, Forrester Cole, Tali Dekel, Andrew Zisserman, William T Freeman, and Michael Rubinstein. Omnimatte: Associating objects and their effects in video. In *CVPR*, 2021.
- [33] Ivana Mikic, Pamela C Cosman, Greg T Kogut, and Mohan M Trivedi. Moving shadow and object detection in traffic scenes. In *ICPR*, 2000.
- [34] Vu Nguyen, Tomas F Yago Vicente, Maozheng Zhao, Minh Hoai, and Dimitris Samaras. Shadow detection with conditional generative adversarial networks. In *ICCV*, 2017.
- [35] Liangqiong Qu, Jiandong Tian, Shengfeng He, Yandong Tang, and Rynson Lau. Deshadownet: A multi-context embedding deep network for shadow removal. In *CVPR*, 2017.
- [36] Jiayu Sun, Ke Xu, Youwei Pang, Lihe Zhang, Huchuan Lu, Gerhard Hancke, and Rynson Lau. Adaptive illumination mapping for shadow detection in raw images. In *ICCV*, 2023.
- [37] Mingxing Tan and Quoc Le. Efficientnet: Rethinking model scaling for convolutional neural networks. In *ICML*, 2019.
- [38] Xinhao Tao, Junyan Cao, Yan Hong, and Li Niu. Shadow generation with decomposed mask prediction and attentive shadow filling. In *AAAI*, 2024.
- [39] Jiandong Tian, Xiaojun Qi, Liangqiong Qu, and Yandong Tang. New spectrum ratio properties and features for shadow detection. *Pattern Recognition*, 51, 2016.

- [40] Kushagra Tiwary, Tzofi Klinghoffer, and Ramesh Raskar. Towards learning neural representations from shadows. In *ECCV*, 2022.
- [41] Tomás F Yago Vicente, Le Hou, Chen-Ping Yu, Minh Hoai, and Dimitris Samaras. Large-scale training of shadow detectors with noisily-annotated shadow examples. In *ECCV*, 2016.
- [42] Hongqiu Wang, Wei Wang, Haipeng Zhou, Huihui Xu, Shaozhi Wu, and Lei Zhu. Language-driven interactive shadow detection. In *ACM MM*, 2024.
- [43] Jifeng Wang, Xiang Li, and Jian Yang. Stacked conditional generative adversarial networks for jointly learning shadow detection and shadow removal. In *CVPR*, 2018.
- [44] Tianyu Wang, Xiaowei Hu, Qiong Wang, Pheng-Ann Heng, and Chi-Wing Fu. Instance shadow detection. In *CVPR*, 2020.
- [45] Xinrui Wang, Lanqing Guo, Xiyu Wang, Siyu Huang, and Bihan Wen. Softshadow: Leveraging penumbra-aware soft masks for shadow removal. In *CVPR*, 2024.
- [46] Tai-Pang Wu, Chi-Keung Tang, Michael S Brown, and Heung-Yeung Shum. Natural shadow matting. *ACM TOG*, 2007.
- [47] Jie Xiao, Xueyang Fu, Yurui Zhu, Dong Li, Jie Huang, Kai Zhu, and Zheng-Jun Zha. Homofomer: Homogenized transformer for image shadow removal. In *CVPR*, 2024.
- [48] Jiamin Xu, Zelong Li, Yuxin Zheng, Chenyu Huang, Renshu Gu, Weiwei Xu, and Gang Xu. Omnisr: Shadow removal under direct and indirect lighting. In *AAAI*, 2025.
- [49] Han Yang, Tianyu Wang, Xiaowei Hu, and Chi-Wing Fu. Silt: Shadow-aware iterative label tuning for learning to detect shadows from noisy labels. In *ICCV*, 2023.
- [50] Xiaofeng Zhang, Chaochen Gu, and Shanying Zhu. Sam-helps-shadow: when segment anything model meet shadow removal. *arXiv:2306.06113*, 2023.
- [51] Quanlong Zheng, Xiaotian Qiao, Ying Cao, and Rynson Lau. Distraction-aware shadow detection. In *CVPR*, 2019.
- [52] Boyan Zhou, Quan Cui, Xiu-Shen Wei, and Zhao-Min Chen. Bbn: Bilateral-branch network with cumulative learning for long-tailed visual recognition. In *CVPR*, 2020.
- [53] Jiejie Zhu, Kegan GG Samuel, Syed Z Masood, and Marshall F Tappen. Learning to recognize shadows in monochromatic natural images. In *CVPR*, 2010.
- [54] Lei Zhu, Zijun Deng, Xiaowei Hu, Chi-Wing Fu, Xuemiao Xu, Jing Qin, and Pheng-Ann Heng. Bidirectional feature pyramid network with recurrent attention residual modules for shadow detection. In *ECCV*, 2018.
- [55] Lei Zhu, Ke Xu, Zhanghan Ke, and Rynson Lau. Mitigating intensity bias in shadow detection via feature decomposition and reweighting. In *ICCV*, 2021.
- [56] Yurui Zhu, Xueyang Fu, Chengzhi Cao, Xi Wang, Qibin Sun, and Zheng-Jun Zha. Single image shadow detection via complementary mechanism. In *ACM MM*, 2022.
- [57] Yurui Zhu, Jie Huang, Xueyang Fu, Feng Zhao, Qibin Sun, and Zheng-Jun Zha. Bijective mapping network for shadow removal. In *CVPR*, 2022.

## NeurIPS Paper Checklist

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [Yes]

Justification: We provide a detailed and accurate description of the paper's contributions and scope in both the abstract and the concluding paragraph of the references.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

### 2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The limitations are mentioned in the discussion section.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [NA]

Justification: The paper does not include theoretical results.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

#### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: We provide the network architectures and hyperparameters required for the experiments in both the paper and supplementary materials.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

#### 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: We will release the dataset and code to the community.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: We provide detailed descriptions of all the training and test details, including backbones, hyperparameters, and optimizers.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [NA]

Justification: Our experiment does not contain statistical analysis.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).

- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: We provide the hardware specifications required for the experiments.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research adheres to the NeurIPS Code of Ethics. It addresses potential societal impacts, avoids harmful applications, and ensures fairness and transparency in methodologies.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [NA]

Justification: There is no societal impact of the work performed.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.

- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: No datasets with potential privacy and security risks appeared in the paper.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: All third-party assets (code, data, models) are properly credited with their original sources cited.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, [paperswithcode.com/datasets](https://paperswithcode.com/datasets) has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.

- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

### 13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [NA]

Justification: The paper does not release new assets.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

### 14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification: This study does not involve any crowdsourced experiments or research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

### 15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]

Justification: The paper does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

## 16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required.

Answer: [NA]

Justification: The core method development in this research does not involve LLMs as any important, original, or non-standard components.

Guidelines:

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.

## Appendix

This Appendix provides extended technical details and experimental analyses to support the main paper. We elaborate on the construction of the FSD dataset across diverse scenarios: indoor/outdoor environments, single/multiple light sources, shadow overlaps, and varying illumination intensities. We describe the physics-based rendering pipeline in Blender Cycles, leveraging path tracing to simulate realistic shadow interactions governed by radiometric principles. We elaborate in detail on the calculation process of the ground truth values for the opacity shadow mask and supplement the qualitative experimental results of shadow detection. Additionally, we further analyze and present the qualitative experimental results of shadow removal using the opacity shadow mask.

## A FSD Dataset

### A.1 Additional details

When constructing the dataset, we considered a variety of scenarios encompassing indoor and outdoor environments, varying from single to multiple light sources, weak to strong lighting conditions, and situations with overlapping shadows. The dataset comprises shadow-free images, instance masks, binary shadow masks, self-shadow masks, and opacity shadow masks.

For indoor object models, we source 1030 models from Google Scanned Objects [7]. To prevent interference within the dataset, we manually exclude cases where a model consists of multiple separate parts. Human models are selected from free models on Sketchfab [21] based on quality, resulting in a manual selection of 59 models. In outdoor settings, we incorporate HDRI textures from Ambient CG [4] and Poly Haven. To ensure data quality, scenes with extremely weak lighting or existing shadows are excluded, leading to a final selection of 26 scenes. Below are some examples of scenes included in the dataset.



Figure 8: Indoor, single light

**Indoor, Single Light Scenario.** In indoor settings, a cubic space serves as the base for constructing the scenes. A total of 428 floor textures from Poly Haven are utilized to adorn the indoor space. Subsequently, 1-5 objects are randomly positioned near the world coordinate origin, with their sizes randomly adjusted to diversify the dataset. The placement strategy for objects in this single-light indoor scenario (refer to Figure. 8) involves distributing the xyxy coordinates within a narrow strip area to simulate common real-world object arrangements. For this scenario, a single light source (such as parallel light or SUN in Blender) is employed to replicate indoor reflection effects, aligning with typical indoor scene configurations. The variability in light sources encompasses factors like intensity, blur level (indicating the softness of shadows), light direction, position, and other parameters. Regarding the camera settings in both indoor and outdoor scenes, the position and angle of the camera are randomized, excluding specifications such as focal length and sensor size.



Figure 9: Indoor, multiple lights

**Indoor, multiple lights.** In more intricate scenarios (refer to Figure. 9), the presence of multiple light sources may give rise to various shadow phenomena. This particular scenario has received limited attention in prior research. To enrich the dataset’s scope, we have included this scenario to expand its comprehensiveness. In this setting, objects are placed with completely random xyxy coordinates. As for lighting, we have implemented a strategy to randomly generate lighting effects, encompassing point light (POINT in Blender), sunlight (parallel light, SUN), spotlight (SPOT), and area light (AREA). The selection of light types and quantities is randomized, and the attributes of the lights are determined in a similar random fashion as previously described.

**Shadow overlaps.** When multiple light sources are utilized, shadow intersections occur (see Figure. 9), a phenomenon that cannot be adequately elucidated by a binary shadow mask alone. Our dataset intentionally includes scenarios featuring shadow intersections. The distinction between this and multiple-light scenarios lies in the placement strategy of objects and lights. Specifically, two objects are positioned to form a specific angle, after which two point light sources are placed in the directions of these objects (from the origin to each object) to generate intersecting shadows.



Figure 10: Outdoor, weak light

**Outdoor, weak light.** For all outdoor scenes, we create a hemisphere and utilize HDRI textures to establish the environment, thereby integrating lighting details from the textures organically. In instances of low lighting intensity (refer to Figure. 10), shadows exhibit a blurred and ghosted appearance, rendering binary shadow masks less efficient for shadow elimination. Given that the lighting details are predetermined by the texture, the only elements subject to randomness are the camera angles and the placement of objects.




Figure 11: Outdoor, strong light

**Outdoor, strong light.** In outdoor environments characterized by intense lighting conditions (see Figure. 11), representative of real-world settings, shadows tend to be distinct and deep. Our dataset comprises a substantial quantity of such scenes. The primary contrast between these two outdoor setups lies in the HDRI textures employed.

Once the scene is set up, we employ keyframe animation to produce the initial image, a shadow-free rendition, and all desired annotations for each object. Blender sequentially generates the animation frame by frame, which is later subjected to post-processing. Due to inherent noise in the original opacity shadow mask from Blender Cycles, we begin by applying noise reduction techniques. Subsequently, we derive the binary shadow mask based on the opacity shadow mask. Following the processing of all samples, we compile a JSON file to document our datasets in the format specified by SOBA [44]. Comprehensive details regarding the scene setup will be released.

## B Additional Shadow Detection Results

### B.1 Shadow Rendering Process

Blender Cycles uses the path tracing algorithm, which is based on the following equation:

$$L_o(p, \omega_o) = L_e(p, \omega_o) + \int_{\Omega} f_r(p, \omega_i, \omega_o) L_i(p, \omega_i) n \cdot \omega_i d\omega_i \quad (12)$$

It uses radiometry to describe lighting more accurately, resulting in more realistic images. The term  $L_i(p, \omega_i) n \cdot \omega_i$  describes the radiance of the light coming from solid angle  $\omega_i$  in the environment, contributing to the radiance at point  $p$ . Thus,  $\int_{\Omega} L_i(p, \omega_i) n \cdot \omega_i d\omega_i$  represents the contribution of all light rays in the hemisphere  $\Omega$ , i.e. all light in the environment. The term  $f_r$  determines how much light will be reflected. Although with numerous BSDF, BRDF models available, they only determine the material of the object. The determinant of shadows is the term  $L_i$ . The term  $L_e$  represents the object's self-emission.

The ground truth opacity shadow mask  $\alpha_{gt}$  is calculated from the shadow image  $S$  and the shadow-free image  $F$ . The specific steps are as follows: first, convert  $S$  and  $F$  to the YCbCr color space and extract the Y channel, then compute the ratio of  $Y_S$  to  $Y_F$ , apply a low-pass filter to remove noise, and finally obtain  $\alpha_{gt}$  through thresholding.

$$\alpha_{gt} = \max(t, f(\frac{Y_S}{Y_F})) \quad (13)$$

Here,  $Y_S$  and  $Y_F$  represent the Y channels of the shadow image and the shadow-free image in the YCbCr color space, respectively.  $f$  denotes a low-pass filtering operation with  $\sigma = 0.5$ , and  $t$  is the threshold, which is set to 0.1 in this paper.

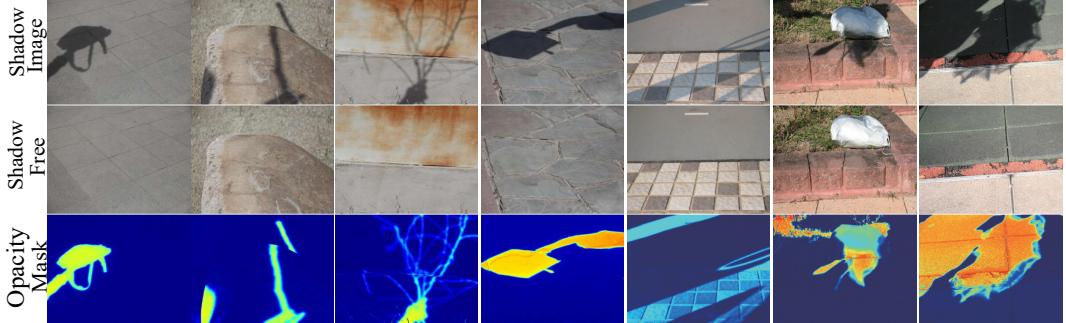


Figure 12: The figure demonstrates the shadow image, the shadow-free image, and the extracted opacity shadow mask. Compared to the binary mask, the opacity shadow mask not only indicates the location of the shadows but also reveals the intensity of the shadows and the degree of background.




Figure 13: The figure illustrates qualitative comparison results between ours and the state-of-the-art methods on the SRD dataset [35], our method exhibits significant advantages in visual quality.

Figure 13 presents additional experimental results on the SRD dataset [35]. Our method outperforms the state-of-the-art approaches across multiple examples, our method demonstrates exceptional performance in fine-grained shadow detection. For instance, in rows 2 and 4, our method can predict the transparency changes of shadows, while in rows 1 and 5, it achieves more precise detection of soft shadow boundaries. These results further validate the reliability of the shadow detection approach guided by shadow opacity.

## B.2 Ablation Study

**Qualitative comparison.** The ablation results in Figure 14 validate our design: the opacity mask enables robust shadow positioning and transparency estimation, complemented by  $\mathcal{L}_{\text{grad}}$  for regulating gradient smoothness and  $\mathcal{L}_{\text{area}}$  for boosting accuracy in extreme scenarios.




Figure 14: Qualitative comparison results of different ablated versions of our model.

**Quantitative Evaluation.** We additionally conducted an ablation study of our loss functions on the ISTD dataset [43]. As it lacks ground truth annotations for shadow opacity, we employed the Balanced Error Rate (BER) metric for evaluation. The results, presented in the table below, demonstrate that our full model outperforms all ablated versions, validating the effectiveness of our design choices.

Method	BER ↓	Shadow↓	Non Shad.↓
w/o $\mathcal{L}_{\text{grad}}$ and $\mathcal{L}_{\text{opa}}$	1.73	<b>0.92</b>	2.53
w/o $\mathcal{L}_{\text{grad}}$ and $\mathcal{L}_{\text{area}}$	1.55	2.21	0.88
w/o $\mathcal{L}_{\text{grad}}$	1.50	2.25	<b>0.75</b>
w/o $\mathcal{L}_{\text{area}}$	<u>1.46</u>	2.15	<u>0.77</u>
Ours (full)	<b>1.32</b>	<u>1.67</u>	0.96

Table 7: Ablation study on ISTD.

## C Additional Shadow Removal Results




Figure 15: We compared the results of using binary masks and opacity masks for shadow removal in training the ShadowDiffusion[12] model on SRD Dataset [26]. The results indicate that using opacity masks instead of binary masks for annotation leads to better visual effects.

In our shadow removal experiments, we improved upon the state-of-the-art ShadowDiffusion [12] method. Specifically, we replaced the binary shadow mask with the detection results from our method for training, while keeping other hyperparameters and optimization strategies consistent with the original ShadowDiffusion. The experiments were conducted on an RTX 4090 GPU, with the training epochs set to 1000. We employed the Adam optimizer (with momentum parameters of  $(0.9, 0.999)$ ) and an initial learning rate of  $3 \times 10^{-5}$ . The model weights were initialized using the Kaiming




Figure 16: Results of different inputs on the LRSS [10] dataset using ShadowDiffusion [12].

initialization technique [15], and an exponential moving average (EMA) strategy with a decay rate of 0.9999 was applied across all experiments.

Figure 15 showcases the performance superiority of the model trained using opacity shadow masks compared to the model trained with binary shadow masks across various test images. The qualitative analysis demonstrates that, under the same training conditions, the utilization of opacity shadow masks enables the model to effectively eliminate shadows while maintaining more detailed information about the underlying scene. Particularly in complex scenarios with blurred shadow boundaries, the model trained with opacity shadow masks shows notably enhanced performance.

To validate the effectiveness of the opacity shadow mask, we conducted supplementary experiments on the LRSS soft shadow dataset [10], which contains 46 pairs of shadow and shadow-free images. The experiments adopted a strategy of training on the SRD dataset [26] and evaluating on the LRSS dataset [10]. As illustrated in Figure. 16, the opacity shadow mask demonstrated significant advantages in restoring soft shadow edges and background textures, thereby fully validating its effectiveness.